

Introduction to nueral network classification with Tensorflow

A classificatoin is where you try to classify something as one thing or another

A few type of classification problems:

- Binary classification
- Multiclass classification
- Multilabel classification

Creating data to view and fit the

```
In [1]: from sklearn.datasets import make_circles # it's the toy dataset

# make 1000 examples
n_samples=1000

# Create cirle
x,y = make_circles(n_samples,
                    noise=0.03,
                    random_state=42)
```

```
In [2]: # Create the features
x[:5]
```

```
Out[2]: array([[ 0.75424625,  0.23148074],
   [-0.75615888,  0.15325888],
   [-0.81539193,  0.17328203],
   [-0.39373073,  0.69288277],
   [ 0.44220765, -0.89672343]])
```

```
In [3]: # create the Labels
y[:5]
```

```
Out[3]: array([1, 1, 1, 1, 0], dtype=int64)
```

as our data is little hard to understand right now... lets visualize it

```
In [4]: import pandas as pd
circles = pd.DataFrame({"x0":x[:,0],"x1":x[:,1],"label":y})
circles
```

	x0	x1	label
0	0.754246	0.231481	1
1	-0.756159	0.153259	1
2	-0.815392	0.173282	1
3	-0.393731	0.692883	1
4	0.442208	-0.896723	0
...
995	0.244054	0.944125	0
996	-0.978655	-0.272373	0
997	-0.136900	-0.810012	1
998	0.670362	-0.767502	0
999	0.281057	0.963824	0

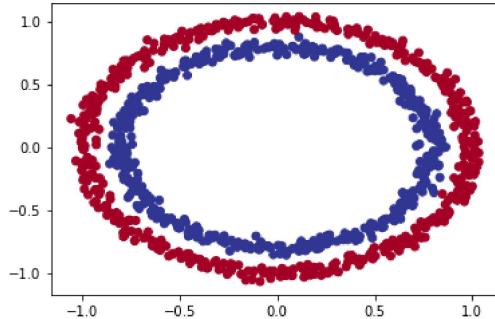
1000 rows × 3 columns

visualize it with plot

```
In [5]: import matplotlib.pyplot as plt
import tensorflow as tf
plt.scatter(x[:,0],x[:,1],c=y,cmap=plt.cm.RdYlBu) # cmap is color dont think its too complex
# also if you have confusion what is x[:,0] then it is this [first_row:last_row,column=0]

# here in this model we try to classify whether our next 1000 data will be blue or red
```

Out[5]: <matplotlib.collections.PathCollection at 0x20297b62fa0>



Input and output shapes

```
In [6]: # check the shapes of our features and Labels
x.shape , y.shape
```

Out[6]: ((1000, 2), (1000,))

```
In [7]: # View tge first example of feature and Labels
x[0],y[0]
```

Out[7]: (array([0.75424625, 0.23148074]), 1)

Steps in modelling

```
In [8]: # Create the model (specified to your problem)
model = tf.keras.Sequential([
    tf.keras.layers.Dense(1)
])

model.compile(loss=tf.keras.losses.BinaryCrossentropy(),
              optimizer=tf.keras.optimizers.SGD(),
              metrics=[ 'accuracy'])

# Fit the model
model.fit(x,y,epochs=100,verbose=0)
```

Out[8]: <keras.callbacks.History at 0x2029e2db850>

```
In [9]: # Evaluate the model
model.evaluate(x,y)
```

32/32 [=====] - 0s 1ms/step - loss: 0.6934 - accuracy: 0.5120

Out[9]: [0.693356990814209, 0.5120000243186951]

Imporoving the model

- Add more layer
- try different optimizers like adam
- try increase the epochs

Imporoved model

```
In [10]: # set the random seed
tf.random.set_seed(42)

# Create the model
model_2 = tf.keras.Sequential([
    tf.keras.layers.Dense(100),
    tf.keras.layers.Dense(10),
    tf.keras.layers.Dense(1)
])

# compile the model
model_2.compile(loss=tf.keras.losses.BinaryCrossentropy(),
                 optimizer=tf.keras.optimizers.SGD(),
                 metrics=['accuracy'])

# Fit the model
model_2.fit(x,y,epochs=100,verbose=0) # fit for 100 passes of the data
```

Out[10]: <keras.callbacks.History at 0x202b1b6e910>

```
In [11]: model_2.evaluate(x,y)
```

32/32 [=====] - 0s 1ms/step - loss: 0.6989 - accuracy: 0.5000

Out[11]: [0.6989450454711914, 0.5]

```
In [12]: model_2.predict(x)
```

32/32 [=====] - 0s 615us/step

```
Out[12]: array([[0.44217896],
                [0.4613756 ],
                [0.4625526 ],
                [0.46633777],
                [0.42571577],
                [0.46720928],
                [0.4631748 ],
                [0.440405 ],
                [0.4359669 ],
                [0.4686557 ],
                [0.44701174],
                [0.45596975],
                [0.46511206],
                [0.43208116],
                [0.42923158],
                [0.43160716],
                [0.4547208 ],
                [0.43219792],
```

```
In [13]: # Set random seed
tf.random.set_seed(42)
```

```
# 1. Create the model (this time 3 Layers)
model_3 = tf.keras.Sequential([
    # Before TensorFlow 2.7.0
    # tf.keras.layers.Dense(100), # add 100 dense neurons

    # With TensorFlow 2.7.0
    # tf.keras.layers.Dense(100, input_shape=(None, 1)), # add 100 dense neurons

    ## After TensorFlow 2.8.0 ##
    tf.keras.layers.Dense(100), # add 100 dense neurons
    tf.keras.layers.Dense(10), # add another layer with 10 neurons
    tf.keras.layers.Dense(1)
])

# 2. Compile the model
model_3.compile(loss=tf.keras.losses.BinaryCrossentropy(),
                 optimizer=tf.keras.optimizers.Adam(), # use Adam instead of SGD
                 metrics=['accuracy'])

# 3. Fit the model
model_3.fit(x, y, epochs=100, verbose=0) # fit for 100 passes of the data
```

Out[13]: <keras.callbacks.History at 0x203bbe69730>

```
In [ ]:
```

Let's visualize it , to visualize it lets create a function plot_decision_boundary

- Takes in trained model , features(x) and labels(y)
- Create a meshgrid of the different x values

- Makes predictions across the meshgrid'
- Plot the prediction as well as a line b/w zones (where each unique class falls)

In [14]: `import numpy as np`

In [15]: `import numpy as np`

```
def plot_decision_boundary(model, x, y):
    """
    Plots the decision boundary created by a model predicting on X.
    This function has been adapted from two phenomenal resources:
    1. CS231n - https://cs231n.github.io/neural-networks-case-study/
    2. Made with ML basics - https://github.com/GokuMohandas/MadeWithML/blob/main/notebooks/08_Neural_Networks.ipynb
    """
    # Define the axis boundaries of the plot and create a meshgrid
    x_min, x_max = x[:, 0].min() - 0.1, x[:, 0].max() + 0.1 # we minus/add with 0.1 to get ourself lil bit some margin
    # it sets up small border b/w the min and max values of each samples so data point are not on very edge
    y_min, y_max = x[:, 1].min() - 0.1, x[:, 1].max() + 0.1
    xx, yy = np.meshgrid(np.linspace(x_min, x_max, 100),
                         np.linspace(y_min, y_max, 100))

    # Create X values (we're going to predict on all of these)
    x_in = np.c_[xx.ravel(), yy.ravel()] # stack 2D arrays together: https://numpy.org/devdocs/reference/generated/numpy.c_.html

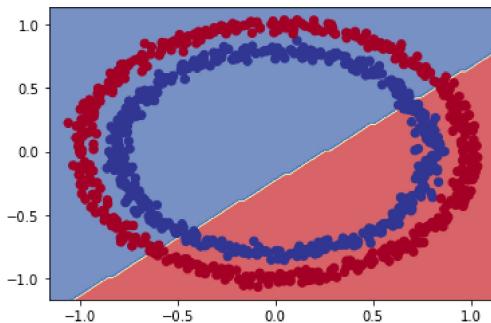
    # Make predictions using the trained model
    y_pred = model.predict(x_in)

    # Check for multi-class
    if len(y_pred[0]) > 1: # checks the final dimension of the model's output shape, if this is > (greater than) 1, it's multi-class
        print("doing multiclass classification...")
        # We have to reshape our predictions to get them ready for plotting
        y_pred = np.argmax(y_pred, axis=1).reshape(xx.shape)
    else:
        print("doing binary classification...")
        y_pred = np.round(np.max(y_pred, axis=1)).reshape(xx.shape)

    # Plot decision boundary
    plt.contourf(xx, yy, y_pred, cmap=plt.cm.RdYlBu, alpha=0.7)
    plt.scatter(x[:, 0], x[:, 1], c=y, s=40, cmap=plt.cm.RdYlBu)
    plt.xlim(xx.min(), xx.max())
    plt.ylim(yy.min(), yy.max())
```

In [16]: `plot_decision_boundary(model_3,x,y)`

313/313 [=====] - 0s 654us/step
doing binary classification...



Let see if our model can be used for a regression problem...

In [17]: `tf.random.set_seed(42)`

```
# Create some regression data
x_regression = tf.range(0,1000,5)
y_regression = tf.range(100,1100,5) # y= x+100

x_regression , y_regression
```

Out[17]: (`<tf.Tensor: shape=(200,), dtype=int32, numpy=`
`array([0, 5, 10, 15, 20, 25, 30, 35, 40, 45, 50, 55, 60,`
 `65, 70, 75, 80, 85, 90, 95, 100, 105, 110, 115, 120, 125,`
 `130, 135, 140, 145, 150, 155, 160, 165, 170, 175, 180, 185, 190,`
 `195, 200, 205, 210, 215, 220, 225, 230, 235, 240, 245, 250, 255,`
 `260, 265, 270, 275, 280, 285, 290, 295, 300, 305, 310, 315, 320,`
 `325, 330, 335, 340, 345, 350, 355, 360, 365, 370, 375, 380, 385,`
 `390, 395, 400, 405, 410, 415, 420, 425, 430, 435, 440, 445, 450,`
 `455, 460, 465, 470, 475, 480, 485, 490, 495, 500, 505, 510, 515,`
 `520, 525, 530, 535, 540, 545, 550, 555, 560, 565, 570, 575, 580,`
 `585, 590, 595, 600, 605, 610, 615, 620, 625, 630, 635, 640, 645,`
 `650, 655, 660, 665, 670, 675, 680, 685, 690, 695, 700, 705, 710,`
 `715, 720, 725, 730, 735, 740, 745, 750, 755, 760, 765, 770, 775,`
 `780, 785, 790, 795, 800, 805, 810, 815, 820, 825, 830, 835, 840,`
 `845, 850, 855, 860, 865, 870, 875, 880, 885, 890, 895, 900, 905,`
 `910, 915, 920, 925, 930, 935, 940, 945, 950, 955, 960, 965, 970,`
 `975, 980, 985, 990, 995])>,`
`<tf.Tensor: shape=(200,), dtype=int32, numpy=`
`array([100, 105, 110, 115, 120, 125, 130, 135, 140, 145, 150,`
 `155, 160, 165, 170, 175, 180, 185, 190, 195, 200, 205,`
 `210, 215, 220, 225, 230, 235, 240, 245, 250, 255, 260,`
 `265, 270, 275, 280, 285, 290, 295, 300, 305, 310, 315,`
 `320, 325, 330, 335, 340, 345, 350, 355, 360, 365, 370,`
 `375, 380, 385, 390, 395, 400, 405, 410, 415, 420, 425,`
 `430, 435, 440, 445, 450, 455, 460, 465, 470, 475, 480,`
 `485, 490, 495, 500, 505, 510, 515, 520, 525, 530, 535,`
 `540, 545, 550, 555, 560, 565, 570, 575, 580, 585, 590,`
 `595, 600, 605, 610, 615, 620, 625, 630, 635, 640, 645,`
 `650, 655, 660, 665, 670, 675, 680, 685, 690, 695, 700,`
 `705, 710, 715, 720, 725, 730, 735, 740, 745, 750, 755,`
 `760, 765, 770, 775, 780, 785, 790, 795, 800, 805, 810,`
 `815, 820, 825, 830, 835, 840, 845, 850, 855, 860, 865,`
 `870, 875, 880, 885, 890, 895, 900, 905, 910, 915, 920,`
 `925, 930, 935, 940, 945, 950, 955, 960, 965, 970, 975,`
 `980, 985, 990, 995, 1000, 1005, 1010, 1015, 1020, 1025, 1030,`
 `1035, 1040, 1045, 1050, 1055, 1060, 1065, 1070, 1075, 1080, 1085,`
 `1090, 1095])>)`

In [18]: `# split into training and testing sets`
`x_reg_train = x_regression[:150]`
`y_reg_train = y_regression[:150]`

`x_reg_test = x_regression[150:]`
`y_reg_test = y_regression[150:]`

```
In [19]: new_model= tf.keras.Sequential([
    tf.keras.layers.Dense(100),
    tf.keras.layers.Dense(10),
    tf.keras.layers.Dense(1)
])

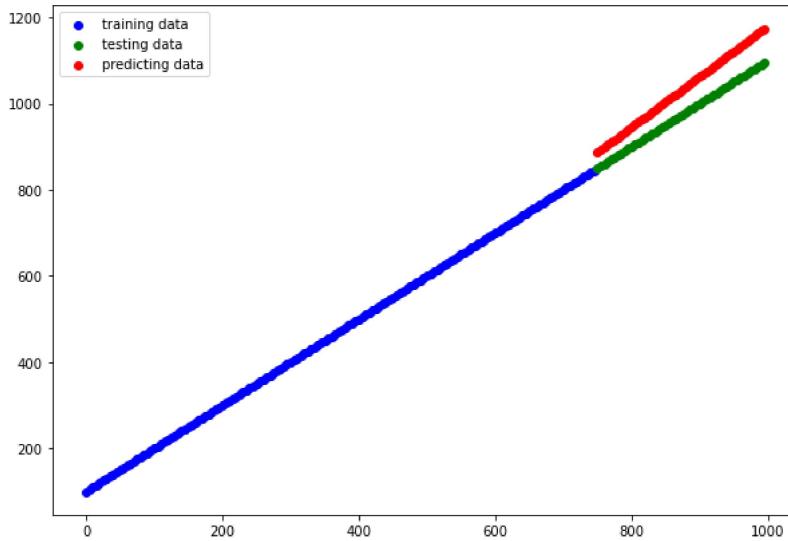
new_model.compile(
    loss=tf.keras.losses.mae,
    optimizer=tf.keras.optimizers.Adam(),
    metrics=["mae"]
)
new_model.fit(tf.expand_dims(x_reg_train, axis=-1),y_reg_train,epochs=100)

Epoch 1/100
5/5 [=====] - 0s 2ms/step - loss: 248.2148 - mae: 248.2148
Epoch 2/100
5/5 [=====] - 0s 1ms/step - loss: 138.8992 - mae: 138.8992
Epoch 3/100
5/5 [=====] - 0s 2ms/step - loss: 53.1035 - mae: 53.1035
Epoch 4/100
5/5 [=====] - 0s 1ms/step - loss: 73.5183 - mae: 73.5183
Epoch 5/100
5/5 [=====] - 0s 2ms/step - loss: 71.2370 - mae: 71.2370
Epoch 6/100
5/5 [=====] - 0s 1ms/step - loss: 47.0045 - mae: 47.0045
Epoch 7/100
5/5 [=====] - 0s 2ms/step - loss: 45.9374 - mae: 45.9374
Epoch 8/100
5/5 [=====] - 0s 1ms/step - loss: 42.3635 - mae: 42.3635
Epoch 9/100
5/5 [=====] - 0s 2ms/step - loss: 43.6835 - mae: 43.6835
Epoch 10/100
5/5 [=====] - 0s 2ms/step - loss: 43.6835 - mae: 43.6835
```

```
In [20]: # Make prediction with our trained model
y_reg_preds = new_model.predict(x_reg_test)

# Plot the model's prediction against our regression model
plt.figure(figsize=(10,7))
plt.scatter(x_reg_train,y_reg_train,c="blue",label="training data")
plt.scatter(x_reg_test,y_reg_test,c="g",label="testing data")
plt.scatter(x_reg_test,y_reg_preds,c="r",label="predicting data")
plt.legend();
```

2/2 [=====] - 0s 989us/step



Common ways to improve deep model:

Adding layers

Increase the numbers of hidden unit:

Change the activation functions

Change the optimization functions

Change the learning rate**Fitting more data****fitting for longer**

THE MISSING PIECE : NON-LINEARITY

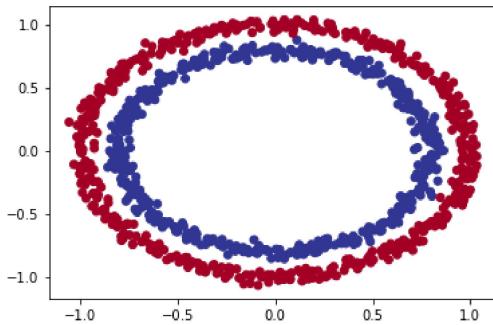
```
In [21]: # Set the random seed
tf.random.set_seed(42)

# 1. Create the model
model_4 = tf.keras.Sequential([
    tf.keras.layers.Dense(1, activation=tf.keras.activations.linear) # here setting the activation function linear
])

# 2. Compile the model
model_4.compile(loss=tf.keras.losses.BinaryCrossentropy(),
                 optimizer=tf.keras.optimizers.Adam(learning_rate=0.001), # note: "lr" used to be what was used, now "Learning_rate"
                 metrics=["accuracy"])
# 3. Fit the model
# history=model_4.fit(tf.expand_dims(x,axis=-1),y,epochs=100) # use tf.expand_dims only when it show error while fitting the mode
model_4.fit(x,y,epochs=100)
```

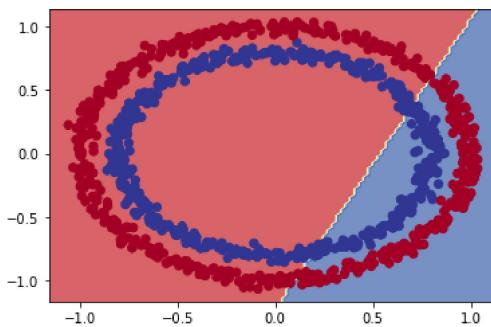
```
Epoch 1/100
32/32 [=====] - 0s 2ms/step - loss: 4.2979 - accuracy: 0.4670
Epoch 2/100
32/32 [=====] - 0s 1ms/step - loss: 4.2317 - accuracy: 0.4400
Epoch 3/100
32/32 [=====] - 0s 2ms/step - loss: 4.1610 - accuracy: 0.4310
Epoch 4/100
32/32 [=====] - 0s 2ms/step - loss: 4.1183 - accuracy: 0.4270
Epoch 5/100
32/32 [=====] - 0s 2ms/step - loss: 4.0784 - accuracy: 0.4240
Epoch 6/100
32/32 [=====] - 0s 2ms/step - loss: 3.9575 - accuracy: 0.4170
Epoch 7/100
32/32 [=====] - 0s 2ms/step - loss: 3.8372 - accuracy: 0.4090
Epoch 8/100
32/32 [=====] - 0s 1ms/step - loss: 3.7656 - accuracy: 0.4120
Epoch 9/100
32/32 [=====] - 0s 1ms/step - loss: 3.7378 - accuracy: 0.4170
Epoch 10/100
32/32 [=====] - 0s 1ms/step - loss: 3.7190 - accuracy: 0.4200
```

```
In [22]: # check out our data
plt.scatter(x[:, 0], x[:, 1], c=y, cmap=plt.cm.RdYlBu);
```



```
In [23]: # Check the decision boundary for our Latest model
plot_decision_boundary(model=model_4,x=x,y=y)
```

313/313 [=====] - 0s 575us/step
doing binary classifcation...



Well, it looks like we're getting a straight (linear) line prediction again.

But our data is non-linear (not a straight line)...

What we're going to have to do is add some non-linearity to our model.

To do so, we'll use the activation parameter in on of our layers.

lets build our nueral network with a non-linear activation function

```
In [24]: # set the random seed
tf.random.set_seed(42)

# 1. Create the model
model_5 = tf.keras.Sequential([
    tf.keras.layers.Dense(1,activation=tf.keras.activations.relu),
    tf.keras.layers.Dense(1)
])

# 2. Compile the model
model_5.compile(loss=tf.keras.losses.binary_crossentropy,
                 optimizer=tf.keras.optimizers.Adam(),
                 metrics=["accuracy"])

# 3. Fit the model
history = model_5.fit(x,y,epochs=100,verbose=0)
```

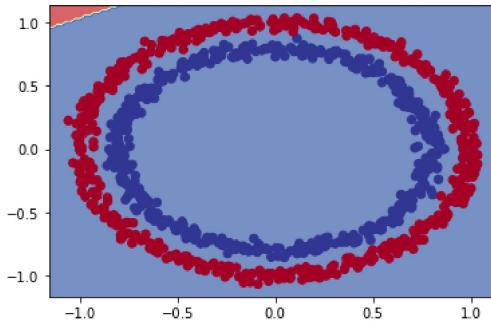
```
In [25]: model_5.evaluate(x,y)
```

32/32 [=====] - 0s 2ms/step - loss: 0.6932 - accuracy: 0.5000

Out[25]: [0.6931846141815186, 0.5]

```
In [26]: plot_decision_boundary(model_5,x,y)
```

313/313 [=====] - 0s 716us/step
doing binary classifcation...



Lets make more better model time to replicate multi-layer nueral network from tensorflow

```
In [27]: # set the random seed
tf.random.set_seed(42)

# 1. Make the model
model_6= tf.keras.Sequential([
    tf.keras.layers.Dense(4,activation="relu"),
    tf.keras.layers.Dense(4,activation="relu")
])

# 2. Compile the model
model_6.compile(loss="binary_crossentropy",
                 optimizer=tf.keras.optimizers.Adam(lr=0.001),
                 metrics=["accuracy"])

# 3. Fit the model
model_6.fit(x,y,epochs=100)
```

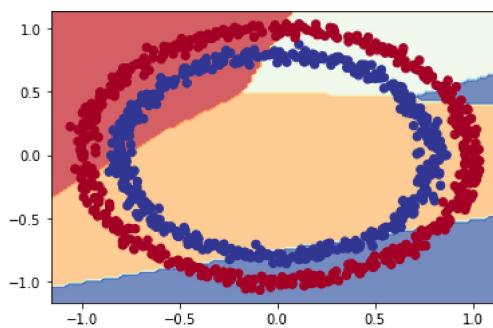
Epoch 1/100

```
C:\Users\Amit\Desktop\sample\anaconda3\lib\site-packages\keras\optimizers\optimizer_v2\adam.py:110: UserWarning: The `lr` argument is deprecated, use `learning_rate` instead.
      super(Adam, self).__init__(name, **kwargs)

32/32 [=====] - 0s 2ms/step - loss: 3.2090 - accuracy: 0.0010
Epoch 2/100
32/32 [=====] - 0s 2ms/step - loss: 2.6774 - accuracy: 0.0050
Epoch 3/100
32/32 [=====] - 0s 2ms/step - loss: 2.5260 - accuracy: 0.0100
Epoch 4/100
32/32 [=====] - 0s 2ms/step - loss: 2.4545 - accuracy: 0.0220
Epoch 5/100
32/32 [=====] - 0s 2ms/step - loss: 2.3271 - accuracy: 0.0420
Epoch 6/100
32/32 [=====] - 0s 1ms/step - loss: 1.9320 - accuracy: 0.0790
Epoch 7/100
32/32 [=====] - 0s 2ms/step - loss: 1.6531 - accuracy: 0.1220
Epoch 8/100
```

```
In [28]: model.evaluate(x,y),
plot_decision_boundary(model_6,x,y)
```

```
32/32 [=====] - 0s 1ms/step - loss: 0.6934 - accuracy: 0.5120
313/313 [=====] - 0s 547us/step
doing multiclass classification...
```



meme material 🤣

This time we'll change the activation function on our output layer too.

Remember the architecture of a classification model? For binary classification, the output layer activation is usually the Sigmoid activation function.

```
In [29]: # set the random seed
tf.random.set_seed(42)

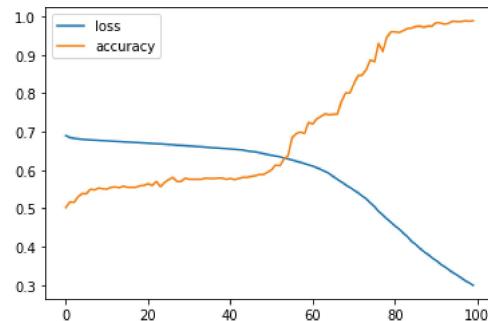
# 1. Make the model
model_6=tf.keras.Sequential([
    tf.keras.layers.Dense(4,activation="relu"), # Layer-1
    tf.keras.layers.Dense(4,activation="relu"), # hidden Layer-2
    tf.keras.layers.Dense(1,activation="sigmoid") # output layer
])

# 2. Compile the model
model_6.compile(loss="binary_crossentropy",
                 optimizer=tf.keras.optimizers.Adam(lr=0.001),
                 metrics=["accuracy"])

# 3. Fit the model
history=model_6.fit(x,y,epochs=100)
```

Epoch 17/100
32/32 [=====] - 0s 2ms/step - loss: 0.6722 - accuracy: 0.5550
Epoch 18/100
32/32 [=====] - 0s 2ms/step - loss: 0.6716 - accuracy: 0.5550
Epoch 19/100
32/32 [=====] - 0s 2ms/step - loss: 0.6710 - accuracy: 0.5590
Epoch 20/100
32/32 [=====] - 0s 2ms/step - loss: 0.6702 - accuracy: 0.5600
Epoch 21/100
32/32 [=====] - 0s 2ms/step - loss: 0.6696 - accuracy: 0.5640
Epoch 22/100
32/32 [=====] - 0s 2ms/step - loss: 0.6688 - accuracy: 0.5600
Epoch 23/100
32/32 [=====] - 0s 2ms/step - loss: 0.6682 - accuracy: 0.5700
Epoch 24/100
32/32 [=====] - 0s 2ms/step - loss: 0.6679 - accuracy: 0.5570
Epoch 25/100
32/32 [=====] - 0s 2ms/step - loss: 0.6671 - accuracy: 0.5670
Epoch 26/100
32/32 [=====] - 0s 2ms/step - loss: 0.6663 - accuracy: 0.5750

```
In [30]: pd.DataFrame(history.history).plot();
```



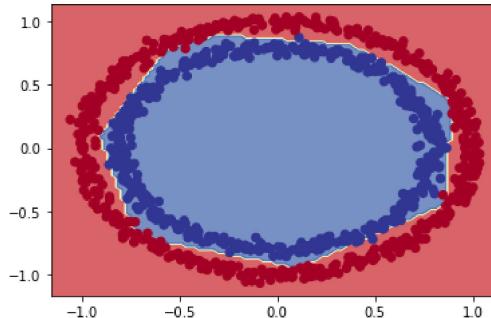
```
In [31]: model_6.evaluate(x,y)
```

```
32/32 [=====] - 0s 1ms/step - loss: 0.2948 - accuracy: 0.9910
```

```
Out[31]: [0.2948004901409149, 0.9909999966621399]
```

```
In [32]: plot_decision_boundary(model_6,x,y)
```

```
313/313 [=====] - 0s 558us/step
doing binary classifcation...
```



see we got the result

our aim was to get non linear line between red and blue circle see we are getting our white line in between

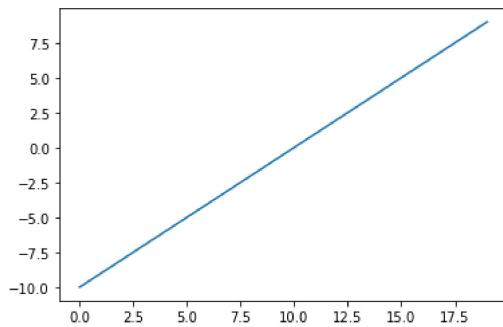
Note: The combination of linear (straight lines) and non-linear (non-straight lines) function is one of the key fundamentals of nueral network

Study about all the activation functions

```
In [33]: # create a toy tensor (similar to the date we pass into our models)
a = tf.cast(tf.range(-10,10),tf.float32)
a
```

```
Out[33]: <tf.Tensor: shape=(20,), dtype=float32, numpy=
array([-10., -9., -8., -7., -6., -5., -4., -3., -2., -1.,  0.,
       1.,  2.,  3.,  4.,  5.,  6.,  7.,  8.,  9.], dtype=float32)>
```

```
In [34]: # visualize our toy tensor
plt.plot(a);
```

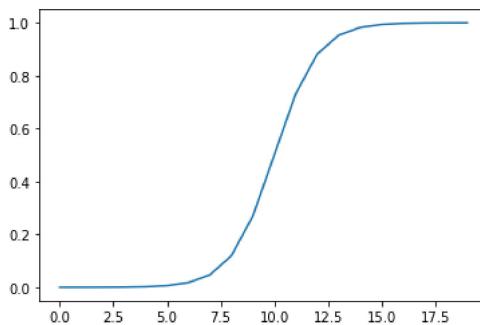


```
In [35]: # Let's start by replicating sigmoid- sigmoid(x) = 1 / (1 + exp(-x)) . non Linear function
def sigmoid (x):
    return 1 / (1 + tf.exp(-x))

# use the sigmoid function on our toy tensor
sigmoid(a)
```

```
Out[35]: <tf.Tensor: shape=(20,), dtype=float32, numpy=
array([4.5397868e-05, 1.2339458e-04, 3.3535014e-04, 9.1105117e-04,
       2.4726230e-03, 6.6928510e-03, 1.7986210e-02, 4.7425874e-02,
       1.1920292e-01, 2.6894143e-01, 5.0000000e-01, 7.3105854e-01,
       8.8079703e-01, 9.5257413e-01, 9.8201376e-01, 9.9330717e-01,
       9.9752742e-01, 9.9908900e-01, 9.9966466e-01, 9.9987662e-01], dtype=float32)>
```

```
In [36]: plt.plot(sigmoid(a));
```

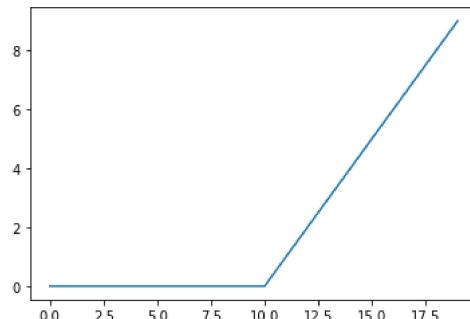


```
In [37]: # use the relu function
# for reference - https://stats.stackexchange.com/questions/226923/why-do-we-use-relu-in-neural-networks-and-how-do-we-use-it
def relu(x):
    return tf.maximum(x,0) # for relu it take only x>0 and other Less than will be converted to 0
relu(a)
```

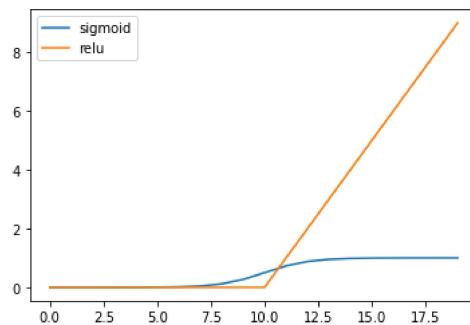
```
Out[37]: <tf.Tensor: shape=(20,), dtype=float32, numpy=
array([0., 0., 0., 0., 0., 0., 0., 0., 0., 1., 2., 3., 4., 5., 6.,
       7., 8., 9.], dtype=float32)>
```

In [38]: `plt.plot(relu(a))`

Out[38]: [`<matplotlib.lines.Line2D at 0x203f70b5100>`]



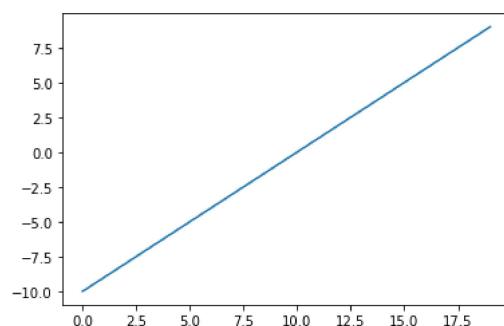
In [39]: `plt.plot(sigmoid(a), label="sigmoid");
plt.plot(relu(a), label="relu");
plt.legend();`



In [40]: `# Let's try the Linear activation function
tf.keras.activations.linear(a)`

Out[40]: `<tf.Tensor: shape=(20,), dtype=float32, numpy=
array([-10., -9., -8., -7., -6., -5., -4., -3., -2., -1., 0.,
 1., 2., 3., 4., 5., 6., 7., 8., 9.],
 dtype=float32)>`

In [41]: `# Does the linear activation function do anything?
plt.plot(tf.keras.activations.linear(a));`



In [42]: `# Does a even changed?
a == tf.keras.activations.linear(a)`

Out[42]: `<tf.Tensor: shape=(20,), dtype=bool, numpy=
array([True, True, True, True, True, True, True, True, True,
 True, True, True, True, True, True, True, True, True,
 True, True])>`

Evaluating and imporoving our classification

So far we've been training and testing on the same dataset...

However, in machine learning this is basically a sin.

So let's create a training and test set.

```
In [43]: # check how many examples we have  
len(x),len(y)
```

```
Out[43]: (1000, 1000)
```

```
In [44]: # split into training and testing sets  
x_train, y_train = x[:800],y[:800]  
x_test, y_test = x[800:],y[800:]  
x_train.shape, y_train.shape , x_test.shape , y_test.shape
```

```
Out[44]: ((800, 2), (800,), (200, 2), (200,))
```

```
In [45]: # Let's recreate a model to fit on the training data and evaluating on the testing dataset

# set the random seed
tf.random.set_seed(42)

# 1 . Creating the model (same as model 7)
model_8 = tf.keras.Sequential([
    tf.keras.layers.Dense(4, activation="relu"),
    tf.keras.layers.Dense(4, activation="relu"),
    tf.keras.layers.Dense(1, activation="sigmoid") # o/p Layer
])

model_8.compile(loss="binary_crossentropy",
                 optimizer=tf.keras.optimizers.Adam(lr=0.01),
                 metrics=["accuracy"])

history=model_8.fit(x_train,y_train,epochs=25)
```

Epoch 1/25

```
C:\Users\Amit\Desktop\sample\anaconda3\lib\site-packages\keras\optimizers\optimizer_v2\adam.py:110: UserWarning: The `lr` argument is deprecated, use `learning_rate` instead.
  super(Adam, self).__init__(name, **kwargs)

25/25 [=====] - 0s 2ms/step - loss: 0.6847 - accuracy: 0.5425
Epoch 2/25
25/25 [=====] - 0s 2ms/step - loss: 0.6777 - accuracy: 0.5525
Epoch 3/25
25/25 [=====] - 0s 1ms/step - loss: 0.6736 - accuracy: 0.5512
Epoch 4/25
25/25 [=====] - 0s 1ms/step - loss: 0.6681 - accuracy: 0.5775
Epoch 5/25
25/25 [=====] - 0s 1ms/step - loss: 0.6633 - accuracy: 0.5850
Epoch 6/25
25/25 [=====] - 0s 2ms/step - loss: 0.6546 - accuracy: 0.5838
Epoch 7/25
25/25 [=====] - 0s 1ms/step - loss: 0.6413 - accuracy: 0.6750
Epoch 8/25
25/25 [=====] - 0s 1ms/step - loss: 0.6264 - accuracy: 0.7013
Epoch 9/25
25/25 [=====] - 0s 1ms/step - loss: 0.6038 - accuracy: 0.7487
Epoch 10/25
25/25 [=====] - 0s 2ms/step - loss: 0.5714 - accuracy: 0.7738
Epoch 11/25
25/25 [=====] - 0s 1ms/step - loss: 0.5404 - accuracy: 0.7650
Epoch 12/25
25/25 [=====] - 0s 2ms/step - loss: 0.5015 - accuracy: 0.7837
Epoch 13/25
25/25 [=====] - 0s 1ms/step - loss: 0.4683 - accuracy: 0.7975
Epoch 14/25
25/25 [=====] - 0s 2ms/step - loss: 0.4113 - accuracy: 0.8450
Epoch 15/25
25/25 [=====] - 0s 992us/step - loss: 0.3625 - accuracy: 0.9125
Epoch 16/25
25/25 [=====] - 0s 1ms/step - loss: 0.3209 - accuracy: 0.9312
Epoch 17/25
25/25 [=====] - 0s 1ms/step - loss: 0.2847 - accuracy: 0.9488
Epoch 18/25
25/25 [=====] - 0s 2ms/step - loss: 0.2597 - accuracy: 0.9525
Epoch 19/25
25/25 [=====] - 0s 1ms/step - loss: 0.2375 - accuracy: 0.9563
Epoch 20/25
25/25 [=====] - 0s 1ms/step - loss: 0.2135 - accuracy: 0.9663
Epoch 21/25
25/25 [=====] - 0s 2ms/step - loss: 0.1938 - accuracy: 0.9775
Epoch 22/25
25/25 [=====] - 0s 1ms/step - loss: 0.1752 - accuracy: 0.9737
Epoch 23/25
25/25 [=====] - 0s 1ms/step - loss: 0.1619 - accuracy: 0.9787
Epoch 24/25
25/25 [=====] - 0s 1ms/step - loss: 0.1550 - accuracy: 0.9775
Epoch 25/25
25/25 [=====] - 0s 2ms/step - loss: 0.1490 - accuracy: 0.9762
```

ValueError: logits and labels must have the same shape ((32, 1) vs (32, 2))

It will give the error when it has more than 1 neuron in output layer in comparison of hidden and input layer

```
In [46]: model_8.evaluate(x_test,y_test)
```

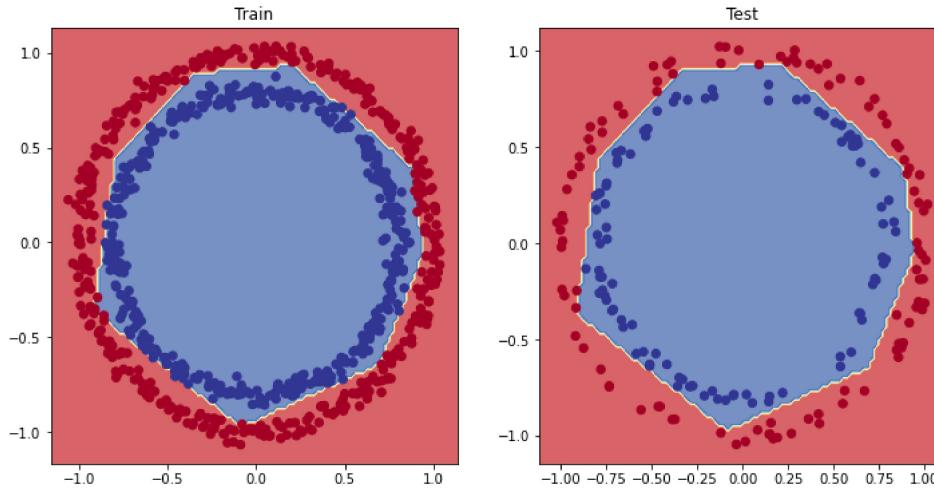
```
7/7 [=====] - 0s 0s/step - loss: 0.1247 - accuracy: 1.0000
```

```
Out[46]: [0.12468849867582321, 1.0]
```

In [47]: # Plot the decision boundaries

```
plt.figure(figsize=(12,6))
plt.subplot(1,2,1) #(row,col,index)
plt.title("Train")
plot_decision_boundary(model_8,x_train,y_train)
plt.subplot(1,2,2)
plt.title("Test")
plot_decision_boundary(model_8,x_test,y_test)
```

```
313/313 [=====] - 0s 627us/step
doing binary classifcation...
313/313 [=====] - 0s 710us/step
doing binary classifcation...
```



In [48]: model_8.summary()

Model: "sequential_8"

Layer (type)	Output Shape	Param #
<hr/>		
dense_18 (Dense)	(None, 4)	12
dense_19 (Dense)	(None, 4)	20
dense_20 (Dense)	(None, 1)	5
<hr/>		
Total params: 37		
Trainable params: 37		
Non-trainable params: 0		

Plot the loss (or training) curves

In [49]: history.history

```
Out[49]: {'loss': [0.6846511960029602,
 0.6777209639549255,
 0.6735945343971252,
 0.6681485772132874,
 0.6632685661315918,
 0.6545671820640564,
 0.6412574648857117,
 0.6264280676841736,
 0.603830873966217,
 0.5714035034179688,
 0.5404428243637085,
 0.5015037059783936,
 0.46833184361457825,
 0.41130149364471436,
 0.36250579357147217,
 0.3209042251110077,
 0.2847079038619995,
 0.2597199082374573,
 0.23746894299983978,
 0.21351952850818634,
 0.19382020831108093,
 0.1752435863018036,
 0.16189301013946533,
 0.1549890637397766,
 0.1489727944135666],
'accuracy': [0.5425000190734863,
 0.5525000095367432,
 0.5512499809265137,
 0.5774999856948853,
 0.5849999785423279,
 0.5837500095367432,
 0.675000011920929,
 0.7012500166893005,
 0.7487499713897705,
 0.7737500071525574,
 0.7649999856948853,
 0.7837499976158142,
 0.7975000143051147,
 0.8450000286102295,
 0.9125000238418579,
 0.9312499761581421,
 0.9487500190734863,
 0.9524999856948853,
 0.956250011920929,
 0.9662500023841858,
 0.9775000214576721,
 0.9737499952316284,
 0.9787499904632568,
 0.9775000214576721,
 0.9762499928474426]}
```

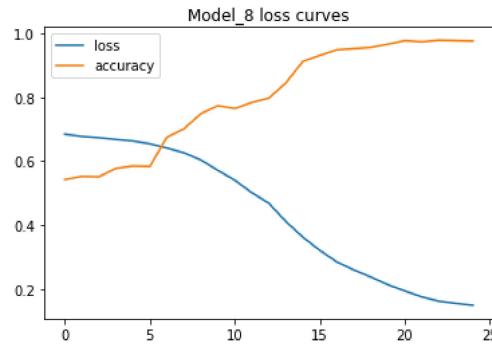
```
In [50]: pd.DataFrame(history.history)
```

```
Out[50]:
```

	loss	accuracy
0	0.684651	0.54250
1	0.677721	0.55250
2	0.673595	0.55125
3	0.668149	0.57750
4	0.663269	0.58500
5	0.654567	0.58375
6	0.641257	0.67500
7	0.626428	0.70125
8	0.603831	0.74875
9	0.571404	0.77375
10	0.540443	0.76500
11	0.501504	0.78375
12	0.468332	0.79750
13	0.411301	0.84500
14	0.362506	0.91250
15	0.320904	0.93125
16	0.284708	0.94875
17	0.259720	0.95250
18	0.237469	0.95625
19	0.213520	0.96625
20	0.193820	0.97750
21	0.175244	0.97375
22	0.161893	0.97875
23	0.154989	0.97750
24	0.148973	0.97625

```
In [51]: # Plot the curves
pd.DataFrame(history.history).plot()
plt.title("Model_8 loss curves")
```

```
Out[51]: Text(0.5, 1.0, 'Model_8 loss curves')
```



Note - For many problems, the loss function going down means the total model is improving (the predictions it's making are getting closer to the ground truth label)

Find the best learning rate

To find the ideal learning rate (the learning rate where the loss decreases the most during training) we're going to use following steps:

- A learning rate **callback** - you can think of a callback as an extra piece of functionality, you can add to your while it's training.
- Another model
- Modified loss curves plot

Type *Markdown* and *LaTeX*: α^2

```
In [52]: # Set random seed
tf.random.set_seed(42)

# Create the model
model_9 = tf.keras.Sequential([
    tf.keras.layers.Dense(4,activation="relu"),
    tf.keras.layers.Dense(4,activation="relu"),
    tf.keras.layers.Dense(1,activation="sigmoid")
])

# compile the model
model_9.compile(loss="binary_crossentropy",
                 optimizer="Adam",
                 metrics=['accuracy'])

# create the learning rate callbacks
lr_schedule = tf.keras.callbacks.LearningRateScheduler(lambda epoch: 1e-4 * 10**(epoch/20)) # its says 10^-4 x 10^(epoch/20)
# # It is a conventional notation for 1 x 10^-4=0.0001

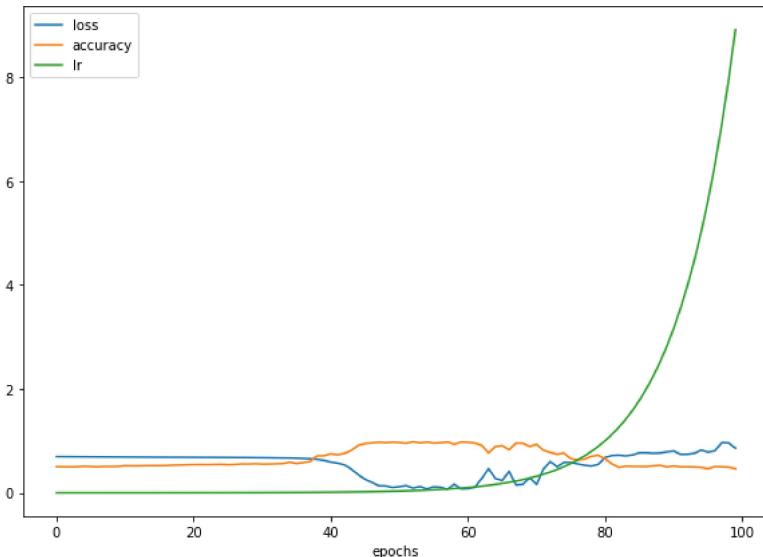
# But it can be ambiguous, as it can also mean (1e)^-4 = 0.018315...
# where e ~2.718281.

# fit the model
history= model_9.fit(x_train,y_train,epochs=100,callbacks=[lr_schedule]) # by the way you can also put more than 1 callbacks
```

```
Epoch 1/100
25/25 [=====] - 0s 2ms/step - loss: 0.6945 - accuracy: 0.4988 - lr: 1.0000e-04
Epoch 2/100
25/25 [=====] - 0s 2ms/step - loss: 0.6938 - accuracy: 0.4975 - lr: 1.1220e-04
Epoch 3/100
25/25 [=====] - 0s 2ms/step - loss: 0.6930 - accuracy: 0.4963 - lr: 1.2589e-04
Epoch 4/100
25/25 [=====] - 0s 1ms/step - loss: 0.6922 - accuracy: 0.4975 - lr: 1.4125e-04
Epoch 5/100
25/25 [=====] - 0s 2ms/step - loss: 0.6914 - accuracy: 0.5063 - lr: 1.5849e-04
Epoch 6/100
25/25 [=====] - 0s 1ms/step - loss: 0.6906 - accuracy: 0.5013 - lr: 1.7783e-04
Epoch 7/100
25/25 [=====] - 0s 2ms/step - loss: 0.6898 - accuracy: 0.4950 - lr: 1.9953e-04
Epoch 8/100
25/25 [=====] - 0s 1ms/step - loss: 0.6889 - accuracy: 0.5038 - lr: 2.2387e-04
Epoch 9/100
25/25 [=====] - 0s 999us/step - loss: 0.6880 - accuracy: 0.5013 - lr: 2.5119e-04
Epoch 10/100
```

```
In [53]: # Check out the history
pd.DataFrame(history.history).plot(figsize=(10,7), xlabel="epochs")
```

```
Out[53]: <AxesSubplot:xlabel='epochs'>
```



In [54]: # Plot the Learning rate vs the Loss

```
lrs = 1e-4 * (10 ** (tf.range(100)/20))
lrs
```

Out[54]: <tf.Tensor: shape=(100,), dtype=float64, numpy=

```
array([1.0000000e-04, 1.12201845e-04, 1.25892541e-04, 1.41253754e-04,
       1.58489319e-04, 1.77827941e-04, 1.99526231e-04, 2.23872114e-04,
       2.51188643e-04, 2.81838293e-04, 3.16227766e-04, 3.54813389e-04,
       3.98107171e-04, 4.46683592e-04, 5.01187234e-04, 5.62341325e-04,
       6.30957344e-04, 7.07945784e-04, 7.94328235e-04, 8.91250938e-04,
       1.0000000e-03, 1.12201845e-03, 1.25892541e-03, 1.41253754e-03,
       1.58489319e-03, 1.77827941e-03, 1.99526231e-03, 2.23872114e-03,
       2.51188643e-03, 2.81838293e-03, 3.16227766e-03, 3.54813389e-03,
       3.98107171e-03, 4.46683592e-03, 5.01187234e-03, 5.62341325e-03,
       6.30957344e-03, 7.07945784e-03, 7.94328235e-03, 8.91250938e-03,
       1.0000000e-02, 1.12201845e-02, 1.25892541e-02, 1.41253754e-02,
       1.58489319e-02, 1.77827941e-02, 1.99526231e-02, 2.23872114e-02,
       2.51188643e-02, 2.81838293e-02, 3.16227766e-02, 3.54813389e-02,
       3.98107171e-02, 4.46683592e-02, 5.01187234e-02, 5.62341325e-02,
       6.30957344e-02, 7.07945784e-02, 7.94328235e-02, 8.91250938e-02,
       1.0000000e-01, 1.12201845e-01, 1.25892541e-01, 1.41253754e-01,
       1.58489319e-01, 1.77827941e-01, 1.99526231e-01, 2.23872114e-01,
       2.51188643e-01, 2.81838293e-01, 3.16227766e-01, 3.54813389e-01,
       3.98107171e-01, 4.46683592e-01, 5.01187234e-01, 5.62341325e-01,
       6.30957344e-01, 7.07945784e-01, 7.94328235e-01, 8.91250938e-01,
       1.0000000e+00, 1.12201845e+00, 1.25892541e+00, 1.41253754e+00,
       1.58489319e+00, 1.77827941e+00, 1.99526231e+00, 2.23872114e+00,
       2.51188643e+00, 2.81838293e+00, 3.16227766e+00, 3.54813389e+00,
       3.98107171e+00, 4.46683592e+00, 5.01187234e+00, 5.62341325e+00,
       6.30957344e+00, 7.07945784e+00, 7.94328235e+00, 8.91250938e+00])>
```

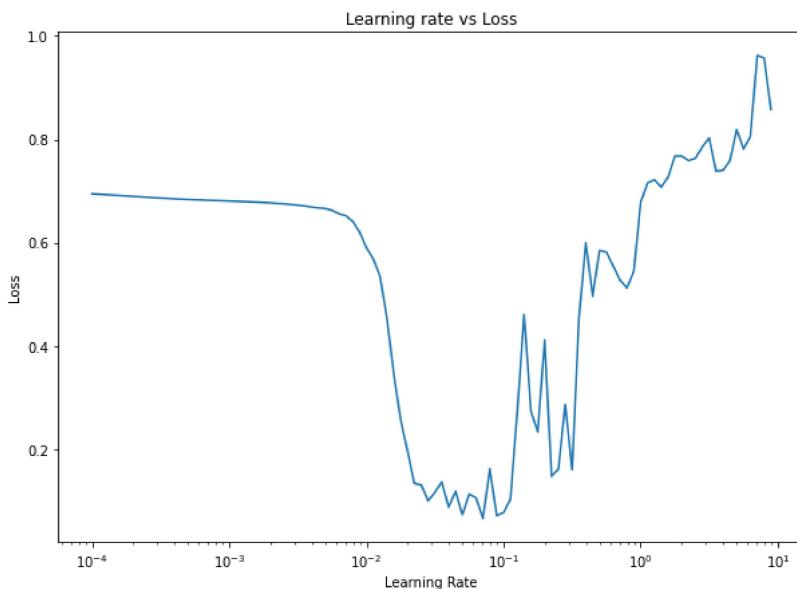
In [55]: len(lrs)

Out[55]: 100

In [56]: # Plot the Learning rate vs the Loss

```
lrs = 1e-4 * (10 ** (tf.range(100)/20))
plt.figure(figsize=(10, 7))
plt.semilogx(lrs,history.history["loss"])
# semilogx( X , Y ) plots x- and y-coordinates using a base-10 logarithmic scale on the x-axis and a Linear scale on the y-axis
plt.xlabel("Learning Rate")
plt.ylabel("Loss")
plt.title("Learning rate vs Loss")
plt.show()
```

ideal learning rate will be at somepoint near it is most lower



```
In [57]: # Example of other typical Learning rate values:  
10**0,10**-1,10**-2,10**-3.1e-4
```

```
Out[57]: (1, 0.1, 0.01, 0.9992864533167718)
```

```
In [58]: # see in the graph the graph comes Lowest at 10**-2 - 2 see b/w 10^-2 and 10^-1 at point 2 its Lowest so it will be 0.02  
10**-2
```

```
Out[58]: 0.01
```

```
In [59]: # Lets train the model with ideal Learning rate  
model_10 = tf.keras.Sequential([  
    tf.keras.layers.Dense(4,activation="relu"),  
    tf.keras.layers.Dense(4,activation="relu"),  
    tf.keras.layers.Dense(1,activation="sigmoid")  
)  
  
model_10.compile(loss="binary_crossentropy",  
                  optimizer=tf.keras.optimizers.Adam(lr=0.01),  
                  metrics=["accuracy"])  
  
# create the Learning rate callbacks  
# lr_schedule = tf.keras.callbacks.LearningRateScheduler(Lambda epoch:1e-4 * 10 **(epoch/20))  
  
history = model_10.fit(x_train, y_train, epochs=25)
```

Epoch 1/25

C:\Users\Amit\Desktop\sample\anaconda3\lib\site-packages\keras\optimizers\optimizer_v2\adam.py:110: UserWarning: The `lr` argument is deprecated, use `learning_rate` instead.
super(Adam, self).__init__(name, **kwargs)

25/25 [=====] - 0s 1ms/step - loss: 0.6925 - accuracy: 0.5188

Epoch 2/25

25/25 [=====] - 0s 1ms/step - loss: 0.6852 - accuracy: 0.5300

Epoch 3/25

25/25 [=====] - 0s 1ms/step - loss: 0.6787 - accuracy: 0.6025

Epoch 4/25

25/25 [=====] - 0s 1ms/step - loss: 0.6662 - accuracy: 0.6538

Epoch 5/25

25/25 [=====] - 0s 2ms/step - loss: 0.6448 - accuracy: 0.6950

Epoch 6/25

25/25 [=====] - 0s 1ms/step - loss: 0.6141 - accuracy: 0.7337

Epoch 7/25

25/25 [=====] - 0s 1ms/step - loss: 0.5652 - accuracy: 0.8487

Epoch 8/25

25/25 [=====] - 0s 1ms/step - loss: 0.5136 - accuracy: 0.8400

Epoch 9/25

25/25 [=====] - 0s 1ms/step - loss: 0.4450 - accuracy: 0.9450

Epoch 10/25

25/25 [=====] - 0s 1ms/step - loss: 0.3721 - accuracy: 0.9712

Epoch 11/25

25/25 [=====] - 0s 1ms/step - loss: 0.3135 - accuracy: 0.9762

Epoch 12/25

25/25 [=====] - 0s 1ms/step - loss: 0.2615 - accuracy: 0.9850

Epoch 13/25

25/25 [=====] - 0s 1ms/step - loss: 0.2167 - accuracy: 0.9850

Epoch 14/25

25/25 [=====] - 0s 1ms/step - loss: 0.1876 - accuracy: 0.9900

Epoch 15/25

25/25 [=====] - 0s 657us/step - loss: 0.1526 - accuracy: 0.9912

Epoch 16/25

25/25 [=====] - 0s 1ms/step - loss: 0.1362 - accuracy: 0.9912

Epoch 17/25

25/25 [=====] - 0s 1ms/step - loss: 0.1141 - accuracy: 0.9937

Epoch 18/25

25/25 [=====] - 0s 1ms/step - loss: 0.1052 - accuracy: 0.9925

Epoch 19/25

25/25 [=====] - 0s 2ms/step - loss: 0.0915 - accuracy: 0.9937

Epoch 20/25

25/25 [=====] - 0s 1ms/step - loss: 0.0830 - accuracy: 0.9962

Epoch 21/25

25/25 [=====] - 0s 1ms/step - loss: 0.0717 - accuracy: 0.9962

Epoch 22/25

25/25 [=====] - 0s 1ms/step - loss: 0.0680 - accuracy: 0.9950

Epoch 23/25

25/25 [=====] - 0s 1ms/step - loss: 0.0595 - accuracy: 0.9987

Epoch 24/25

25/25 [=====] - 0s 1ms/step - loss: 0.0555 - accuracy: 0.9950

Epoch 25/25

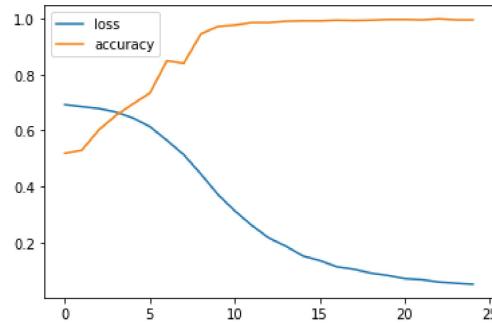
25/25 [=====] - 0s 2ms/step - loss: 0.0517 - accuracy: 0.9950

```
In [60]: model_10.evaluate(x_test,y_test)
```

7/7 [=====] - 0s 0s/step - loss: 0.0579 - accuracy: 0.9850

```
Out[60]: [0.05786607787013054, 0.9850000143051147]
```

```
In [61]: pd.DataFrame(history.history).plot();
```



```
In [62]: model_8.evaluate(x_test,y_test),model_9.evaluate(x_test,y_test),model_10.evaluate(x_test,y_test)
```

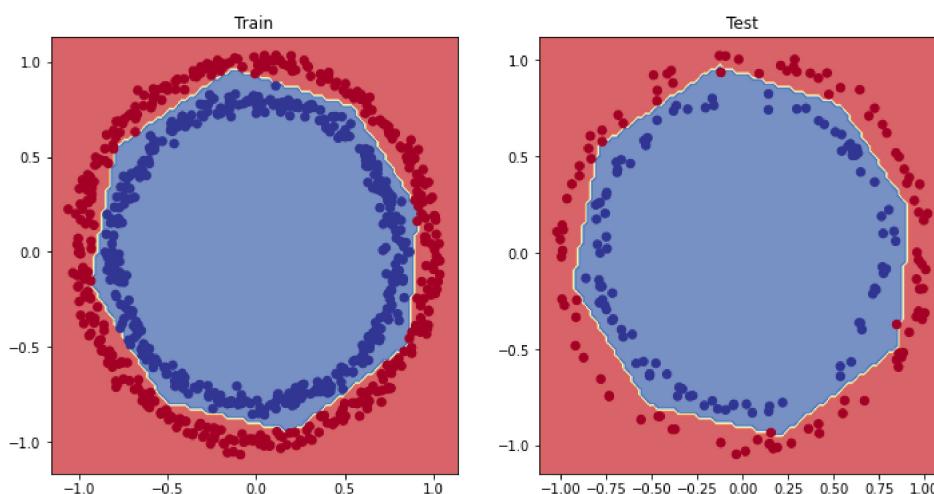
7/7 [=====] - 0s 3ms/step - loss: 0.1247 - accuracy: 1.0000
7/7 [=====] - 0s 3ms/step - loss: 0.6935 - accuracy: 0.4950
7/7 [=====] - 0s 1ms/step - loss: 0.0579 - accuracy: 0.9850

```
Out[62]: ([0.12468849867582321, 1.0],  
 [0.6934590339660645, 0.4950000047683716],  
 [0.05786607787013054, 0.9850000143051147])
```

```
In [63]: # Plot the decision boundaries for the training and test sets
```

```
plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
plt.title("Train")
plot_decision_boundary(model_10, x=x_train, y=y_train)
plt.subplot(1, 2, 2)
plt.title("Test")
plot_decision_boundary(model_10, x=x_test, y=y_test)
plt.show()
```

313/313 [=====] - 0s 554us/step
doing binary classifcation...
313/313 [=====] - 0s 544us/step
doing binary classifcation...



More classification evaluation methods

Alongside visualizing our models results as much as possible, there are a handful of other classification evaluation methods & metrics you should be familiar with:

```
* Accuracy
* Precision
* Recall
* F1-score
* Confusion matrix
* Classification report (from scikit-learn) - https://scikitlearn.org/stable/modules/generated/sklearn.metrics.classification_report.html
```

In [64]:

```
# Check the accuracy of our model
loss, accuracy = model_10.evaluate(x_test,y_test)
print(f"Model loss on the test set: {loss}")
print(f"Model accuracy on the test set:{(accuracy*100):.2f}%")
```

```
7/7 [=====] - 0s 05/step - loss: 0.0579 - accuracy: 0.9850
Model loss on the test set: 0.05786607787013054
Model accuracy on the test set:98.50%
```

How to make confusion matrix?

In [65]:

```
# # Create a confusion matrix
# from sklearn.metrics import confusion_matrix

# # Make predictions
# y_preds = model_10.predict(x_test)
# # confusion matrix
# confusion_matrix(y_test,y_preds)

# give this error
# ValueError: Classification metrics can't handle a mix of binary and continuous targets
```

As you can see Classification metrics can't handle a mix of binary and continuous targets so lets round it up

In [66]:

```
y_test[:5],y_preds[:5]
```

```
-----
NameError                                 Traceback (most recent call last)
Input In [66], in <cell line: 1>()
----> 1 y_test[:5],y_preds[:5]

NameError: name 'y_preds' is not defined
```

In []:

```
# Create a confusion matrix
from sklearn.metrics import confusion_matrix

# Make predictions
y_preds = tf.math.round(model_10.predict(x_test))

# confusion matrix
confusion_matrix(y_test,y_preds)
```

how about we pretify our confusion matrix?

Here created the function so we can check for every model

```
In [ ]: # Note : the confusion matrix code we're about to write is remix of scikit-Learn's plot_confusion_matrix #92

import itertools

def cm_figure(y_true,y_preds):
    figsize = (10,10)

    # Create the confusion matrix
    cm = confusion_matrix(y_true,tf.math.round(y_preds))
    cm_norm = cm.astype("float") / cm.sum(axis=1)[:,np.newaxis] # normalize our confusion matrix
    n_classes = cm.shape[0]

    # Lets prettify it
    fig, ax = plt.subplots(figsize=figsize)
    # Create a matrix plot
    cax = ax.matshow(cm , cmap=plt.cm.Blues) # here cmp cm is colour matrix
    fig.colorbar(cax)

    # Create classes
    classes = False

    if classes:
        labels = classes
    else:
        labels = np.arange(cm.shape[0])

    # Label the axes
    ax.set(title = "Confusion Matrix",
           xlabel= "Predicted Label",
           ylabel= "True Label",
           xticks= np.arange(n_classes),
           yticks= np.arange(n_classes),
           xticklabels=labels,
           yticklabels=labels)

    # Set the x-axis labels to bottom
    ax.xaxis.set_label_position("bottom")
    ax.xaxis.tick_bottom()

    # Adjust the label size
    ax.yaxis.label.set_size(20)
    ax.xaxis.label.set_size(20)
    ax.title.set_size(20)

    # Set the threshold for different color
    threshold = (cm.max() + cm.min()) / 2.

    # Plot the text on each cell
    for i, j in itertools.product(range(cm.shape[0]),range(cm.shape[1])):
        plt.text(j,i,f'{cm[i, j]} ({cm_norm[i, j]*100:.1f}%)',
                 horizontalalignment="center",
                 color="white" if cm[i, j]> threshold else "black",
                 size=15)
```

```
In [ ]: cm_figure(y_test,y_preds)
```

Working with a larger example (multiclass classification)

When you have more than two classes as an option, it's known as multi-class classification.

- This means if you have 3 different classes, it's multi-class classification.
- It also means if you have 100 different classes, it's multi-class classification.

Data-set is:-

To practise multi-class classification, we're going to build a nueral network to classify image of different items of clothing

we are taking tensorflow built-in dataset which is fashion clothing dataset which classify image on basis of clothes, uses 10 classes

```
In [68]: # Lets make different folder for this
```

```
In [ ]:
```

